# Execution architecture and evidence boundary

Researcher-only preparation is separated from the evidence available to each experimental agent

## RESEARCHER-ONLY

Outside every agent workspace

**No agent access to answer key**

### Paired snapshots

Vulnerable / fixed source
Patch and metadata

### Validated scenario card

Threat model and scoped label
Evidence manifest / allowlist

### Scoring reference

Source-based rationale
Researcher judgments

## CONFINED RUN: ONE FRESH HARNESS SESSION

Condition source read-only; isolated runtime home writable

### Initial source packet

Frozen task + condition evidence
No researcher labels or fix

### Condition workspace

One of E1 / E2 / E3
Only allowed vulnerable-source files

### Agent harness

Qwen Code OR OpenCode
Own orchestration / native tool schema

### Read-only tool interface

Read / search / list
Requests, results and errors recorded

### Shared inference service (outside agent process)

Ollama / qwen3.5:9b-q4_K_M / configured context 16,384

OS + tool policy deny private files, sibling conditions and workspace writes.

### External collection and mechanical normalization

Raw stdout / stderr / metadata / final text -> preserved traces and draft CSV

### Researcher review: source + observed trace -> manual scores with citations